



CS3232 Fundamental of Deep learning

Parts of the contents are from deeplearning.ai, TensorFlow, jalFaizy and other resources on the Internet

Generative Adversarial Networks

Generative Adversarial Networks

- Generative Adversarial Networks (GANs) are a class of generative models introduced by Ian Goodfellow and his colleagues in 2014.
- GANs consist of two neural network models, the generator and the discriminator, that are trained together through adversarial processes.
- The goal of GANs is to generate new data samples that are indistinguishable from real data samples.
 - Given a training set, a GAN learns to generate new data with the same statistics as the training set.
 - GANs depend much on the training loss of the model, the model tries to minimise loss to generate as real images as possible.

Generative Adversarial Networks

GENERATOR
"The Artist"
A neural network trying to
create pictures of cats that
look real.



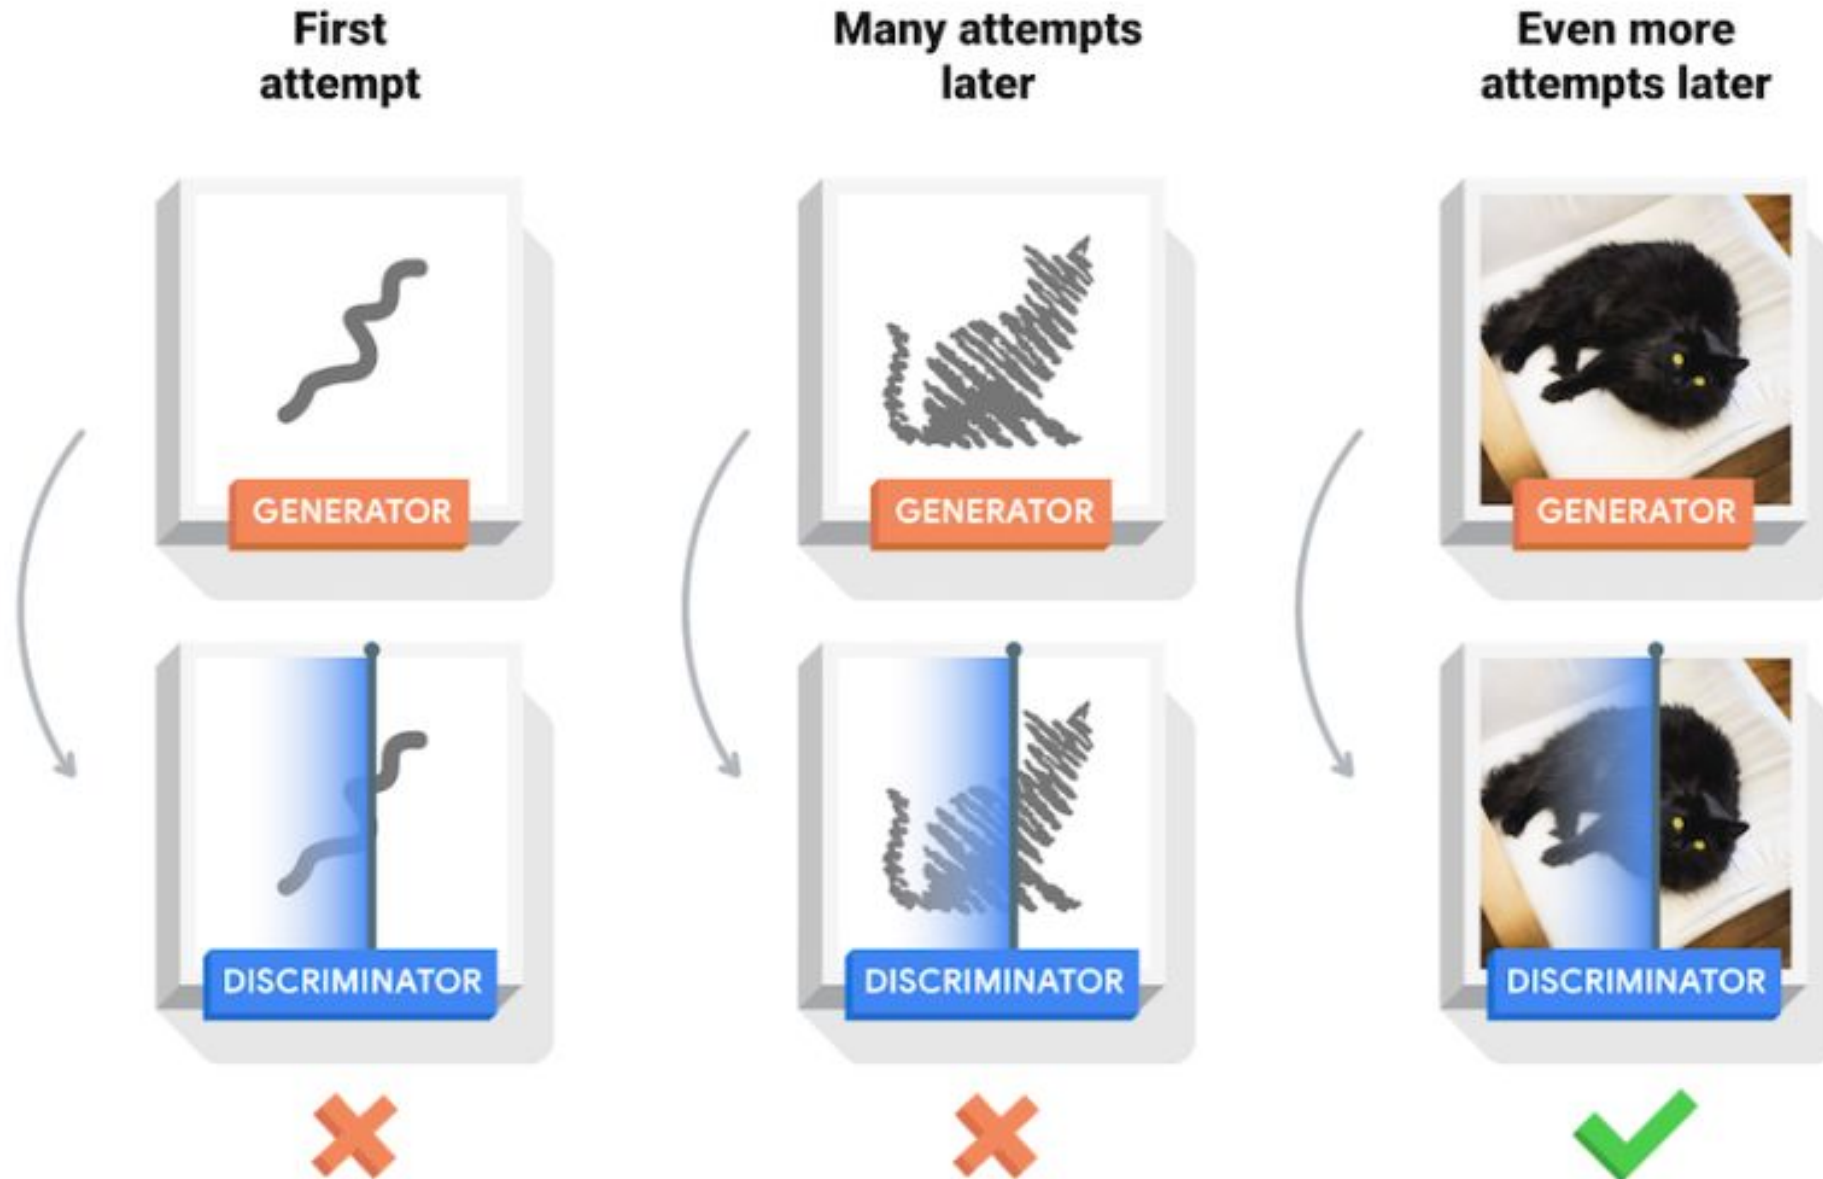
Thousands of real-world
images labeled "CAT"



DISCRIMINATOR
"The Art Critic"
A neural network examining
cat pictures to determine if
they're real or fake.




Generative Adversarial Networks



GAN

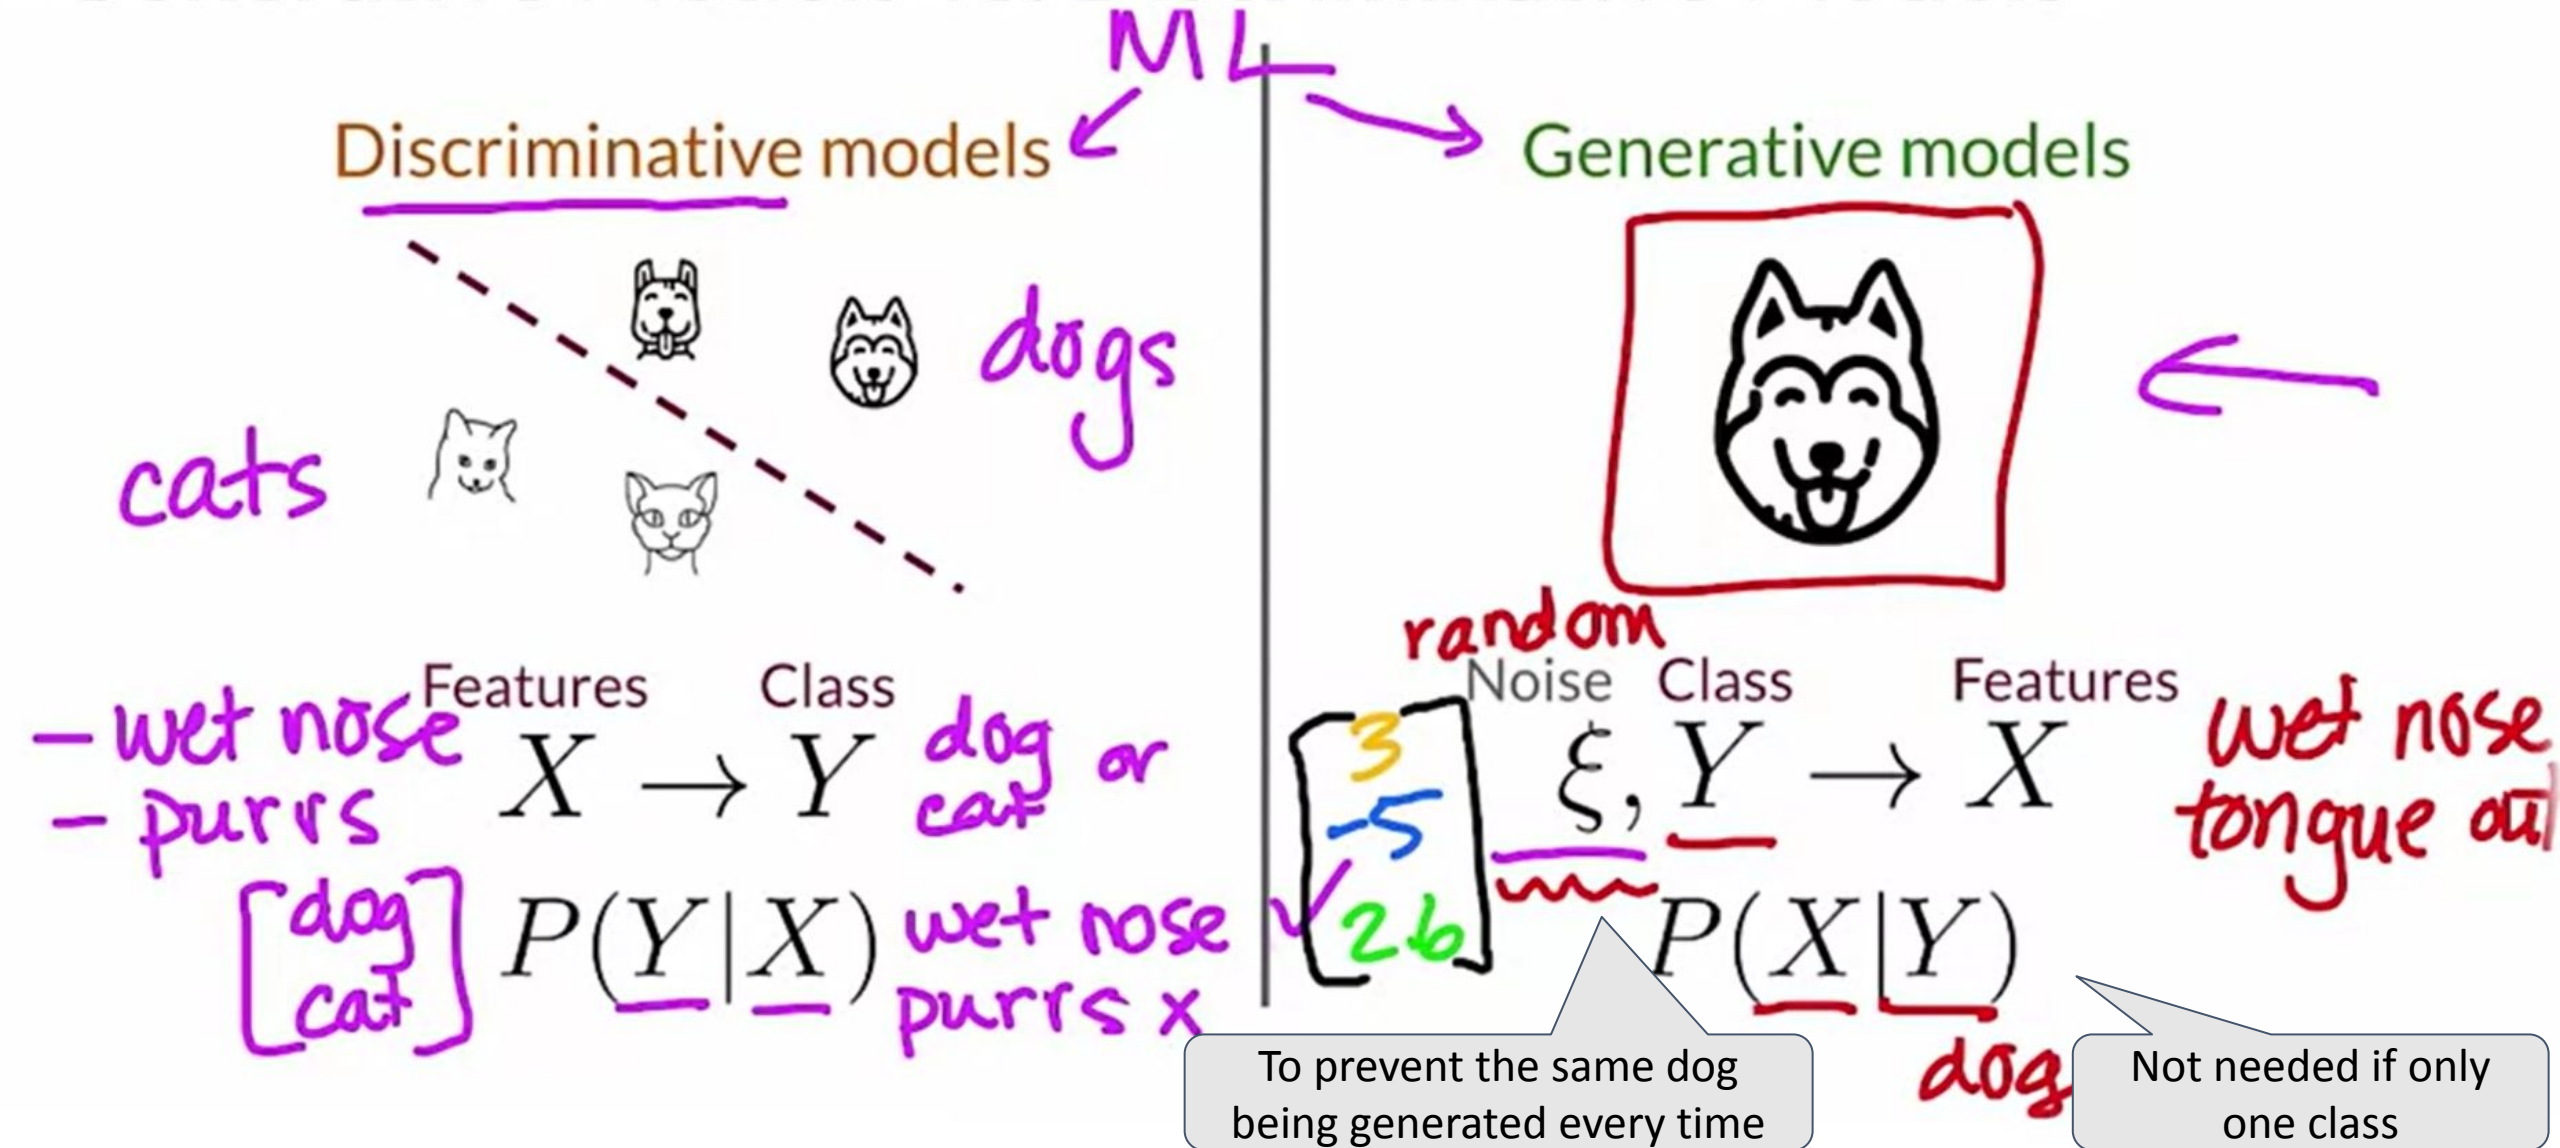
- GAN is a generative model which achieves a high level of realism by pairing a generator with a discriminator.
- The generator learns to produce the target output, while the discriminator learns to distinguish true data from the output of the generator.
 - The generator tries to fool the discriminator, and the discriminator tries to keep from being fooled.
-

Generative Models vs. Discriminative Models



Generative Adversarial Networks


Generative Models vs. Discriminative Models



DMs and GMs mirror each other

Generative Models vs. Discriminative Models

Discriminative models



Features Class

$$X \rightarrow Y$$

$$P(Y|X)$$

Generative models



Noise Class Features

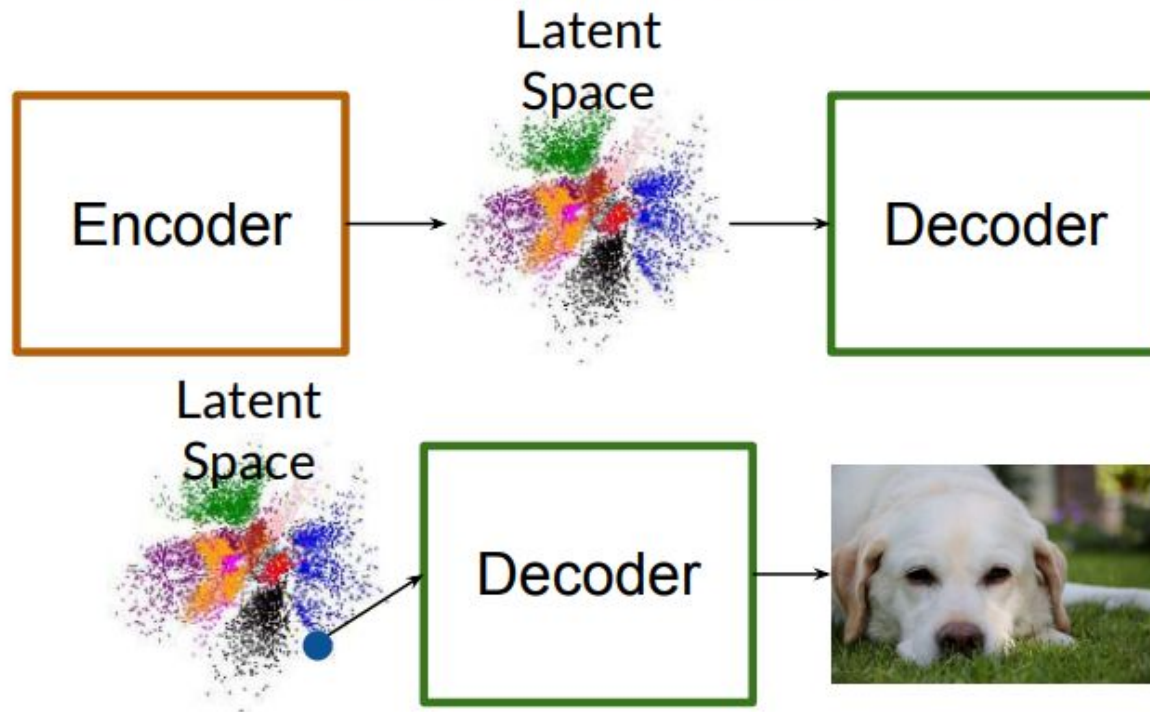
$$\xi, Y \rightarrow X$$

$$P(X|Y)$$

Many types of generative models

Generative Models

Variational Autoencoders



Generative Adversarial Networks

Many types of generative models

Generative Models

Variational Autoencoders

Latent Space

Encoder

Decoder

Latent Space

Decoder



Generative Adversarial Networks

GAN

fake
Compete

Generator

~ decoder

Discriminator

$\begin{bmatrix} 1.2 \\ 3 \\ -5 \end{bmatrix}$

$\begin{bmatrix} -5 \\ 6.2 \\ 8 \end{bmatrix}$

Random Noise

Generator



GANs

- Generative models learn to produce realistic examples
- Discriminative models distinguish between classes
- They both compete in GANs
 - Applications

GANs Over Time



Ian Goodfellow
@goodfellow_ian


Creator of GANs

4.5 years of GAN progress on face generation.

arxiv.org/abs/1406.2661 arxiv.org/abs/1511.06434

arxiv.org/abs/1606.07536 arxiv.org/abs/1710.10196

arxiv.org/abs/1812.04948



GANs Over Time



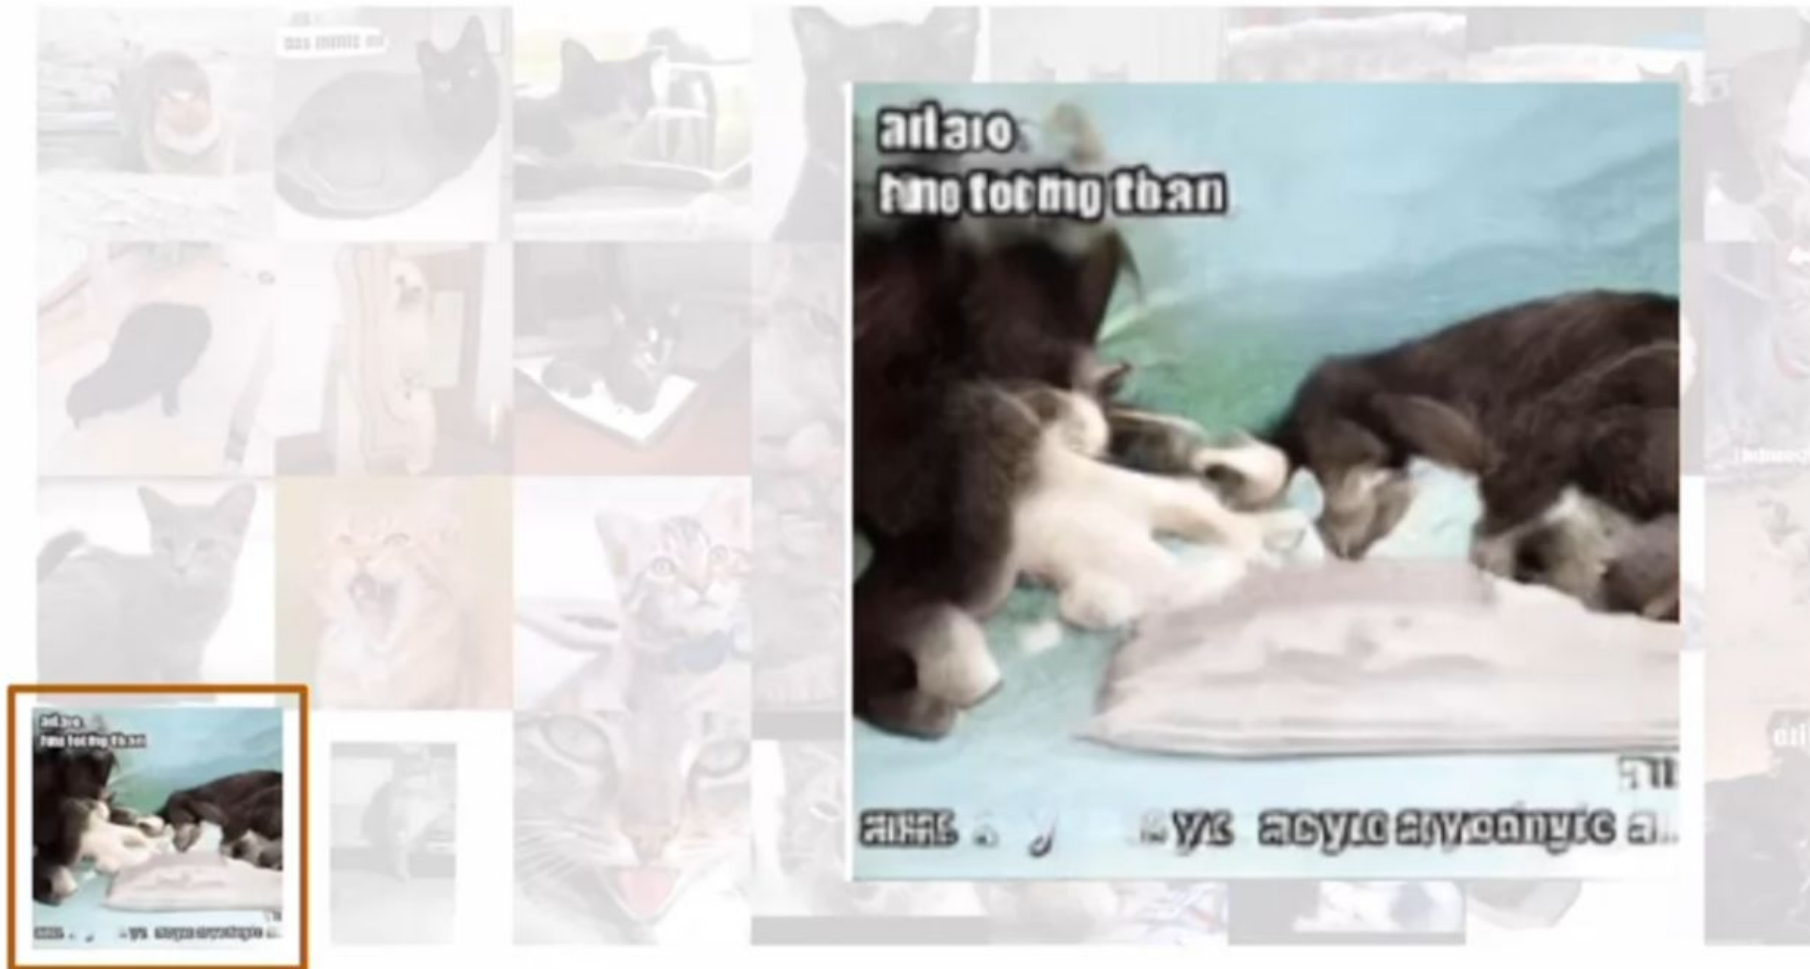
Face Generation
StyleGAN2

These people do
not exist!

Karras, Tero, et al. "Analyzing and improving the image quality of stylegan." *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 2020.

GANs Over Time

Can blur images and have text on them too



StyleGAN2



Mimics the
distribution of
the training data


GANs for Image Translation

From one domain to another

CycleGAN



GANs for Image Translation



GauGAN

Doodles




Pictures

GANs are Magic!

Can animate real life photos



GANs for 3D Objects



Companies Using GANs



Next-gen
Photoshop



Text
Generation



Data
Augmentation



Image Filters



Super-resolution

Intuition behind GANs

- The **generator's** goal is to fool the discriminator
- The **discriminator's** goal is to distinguish between real and fake
- They learn from the competition with each other
- At the end, **fakes** look **real**



Discriminator's goal

Classifiers

Distinguish between different **classes**



Classifier

Turtle

Bird

Cat

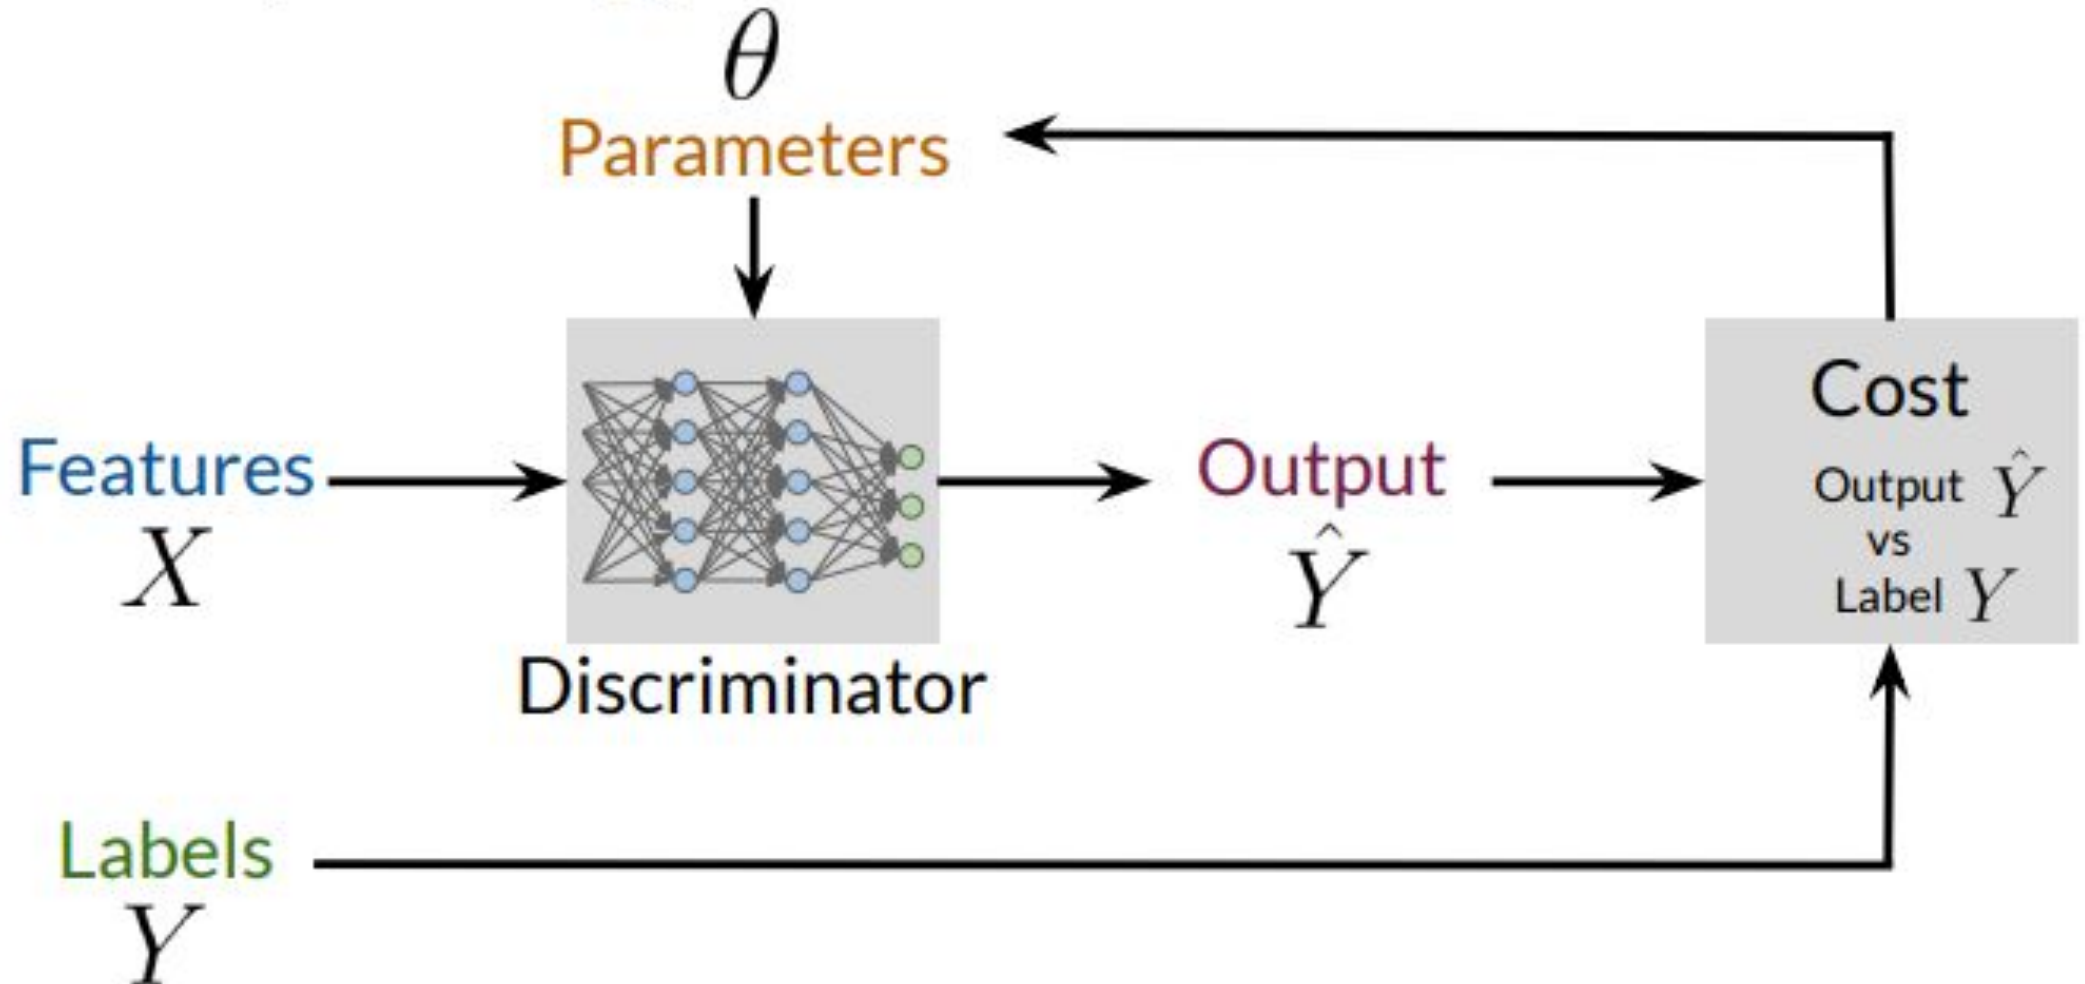
Dog

Fish

IT can classify text also

Discriminator

Classifiers (training)

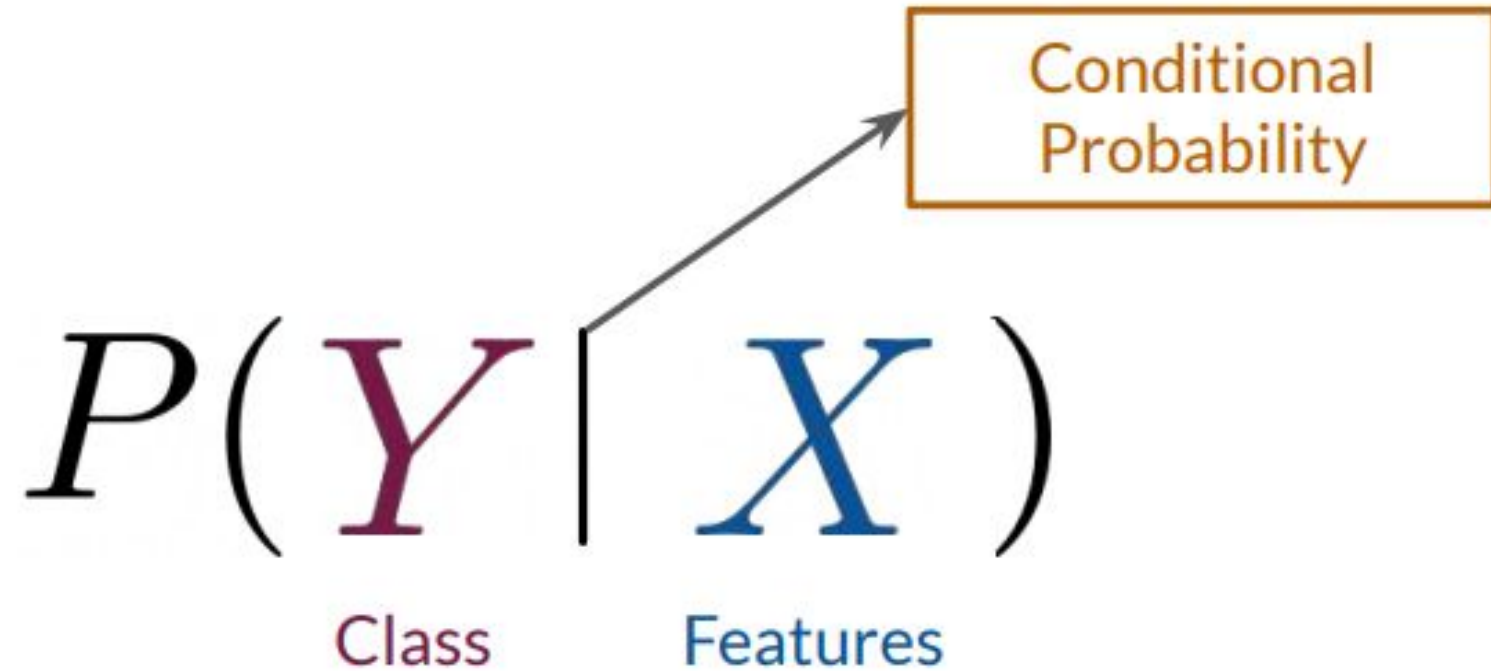


Model probability of each class

Classifiers



$$P\left(\begin{array}{c} \text{Turtle} \\ \text{Bird} \\ \text{Cat} \\ \text{Dog} \\ \text{Fish} \end{array} \mid \text{Image} \right)$$

Classifiers



Classify real vs fake (rather than cat, dog or bird)

Discriminator



Discriminator

$$P(\text{Fake} \mid X)$$

Class Features

Discriminator

$$P\left(\begin{array}{c} \text{Fake} \\ \text{Class} \end{array} \mid \begin{array}{c} \text{Features} \end{array} \right) = 0.85 \rightarrow \boxed{\text{Fake}}$$

Summary

- The **discriminator** is a classifier
- It learns the probability of class Y (**real** or **fake**) given features X
- The probabilities are the feedback for the **generator**

Generator

Turtle

Generates examples of the class

Bird

Cat


Dog

Fish




Generate different examples at different ~~times~~

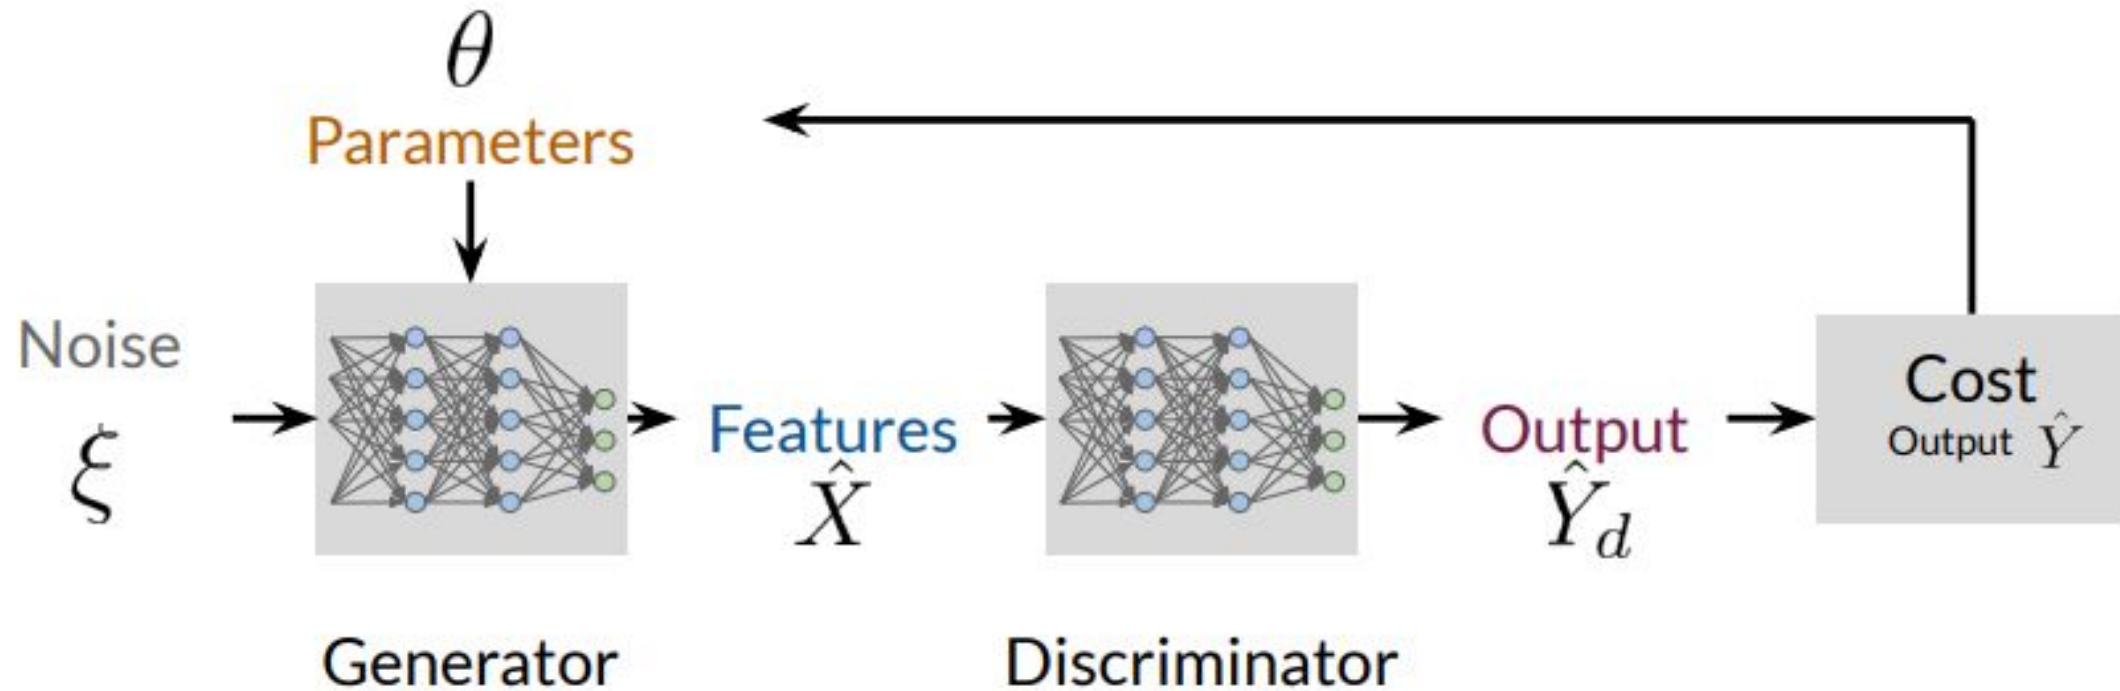
Neural Networks



Neural Networks



Generator: Learning



Sampling



Noise vectors

θ

Saved
Generator



Generator

$$P(\text{)$$

Turtle


Bird

Cat

Dog


Fish

Generator



When only one class, no Y

Generator

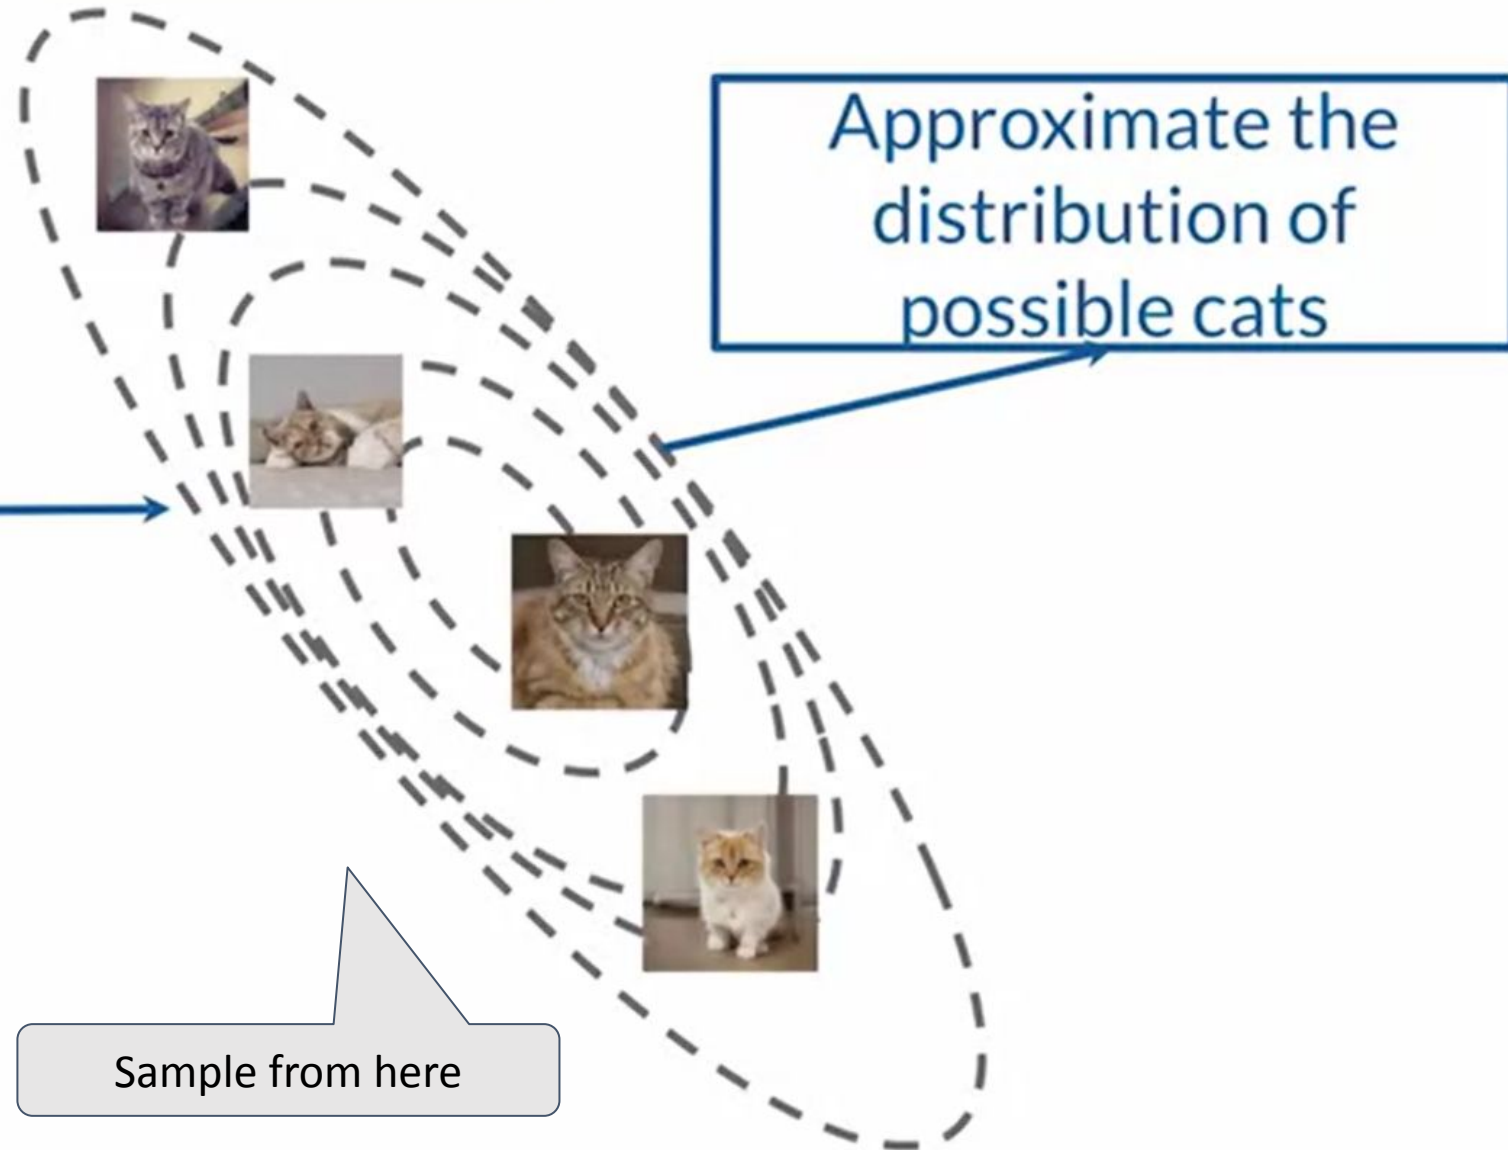


More common types are more likely to be generated

Generator

$P(X)$

Features



Summary

- The **generator** produces fake data
- It learns the probability of features X
- The **generator** takes as input noise (random features)



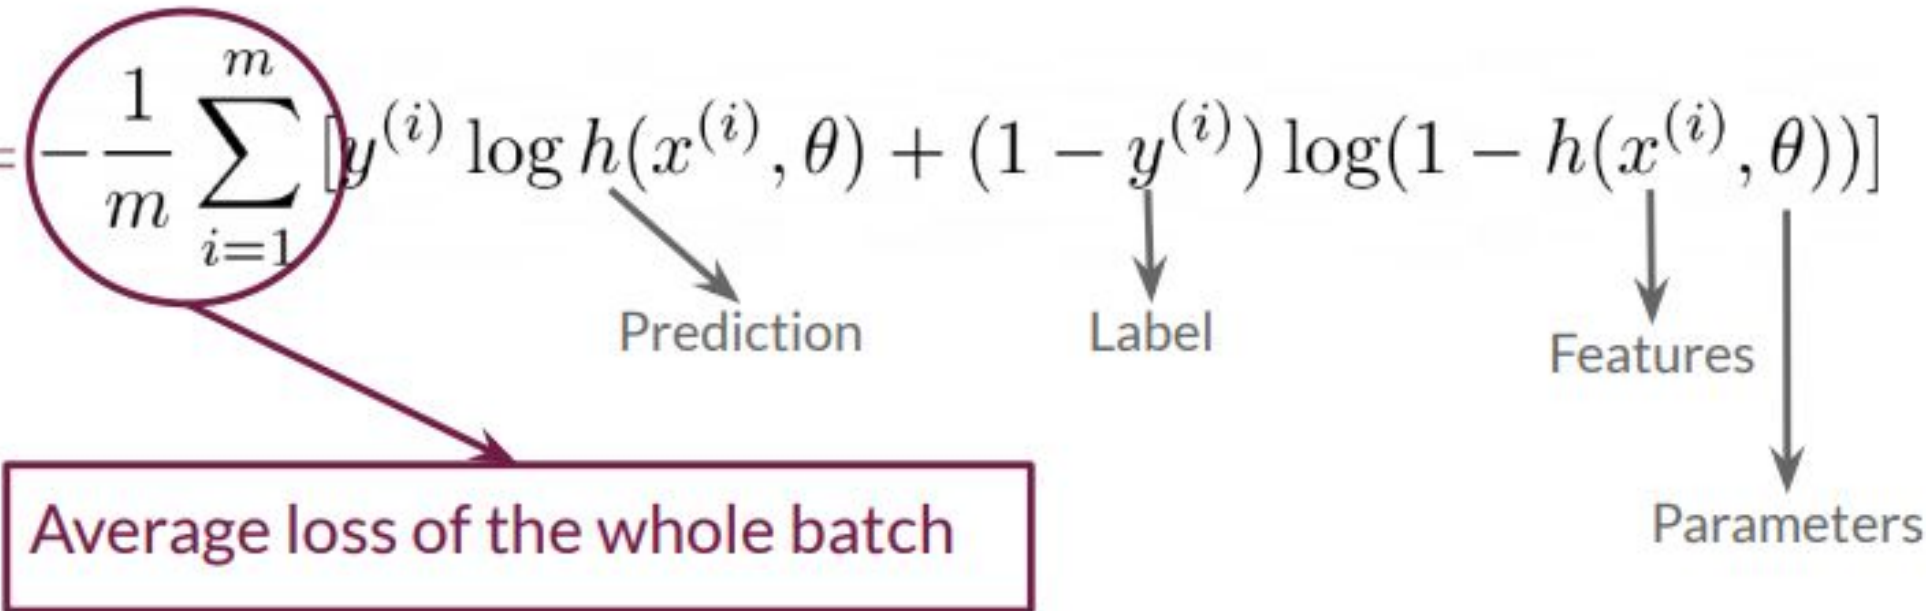
Binary cross entropy function

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$



Average loss of the whole batch

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}; \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

When is this term relevant?
What happens when the prediction is close?
When the prediction is very bad?

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}; \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

Fake

$y^{(i)}$	$h(x^{(i)}, \theta)$	$y^{(i)} \log h(x^{(i)}, \theta)$
0	any	0
1	0.99	~ 0
1	~ 0	$-\text{inf}$

real

Relevant when
the label is 1

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

How about this term?

BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

$y^{(i)}$	$h(x^{(i)}, \theta)$	$(1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))$
1	any	0
0	0.01	~0
0	~1	-inf

Relevant when
the label is 0

BCE Cost Function


$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

Both terms evaluate to $-\infty$ when the predictions are bad which is how it needs to be as a cost function (so that it gets minimized during optimization)

Ensures that the cost is always greater or equal to 0



BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$



BCE Cost Function

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$



Label is 0

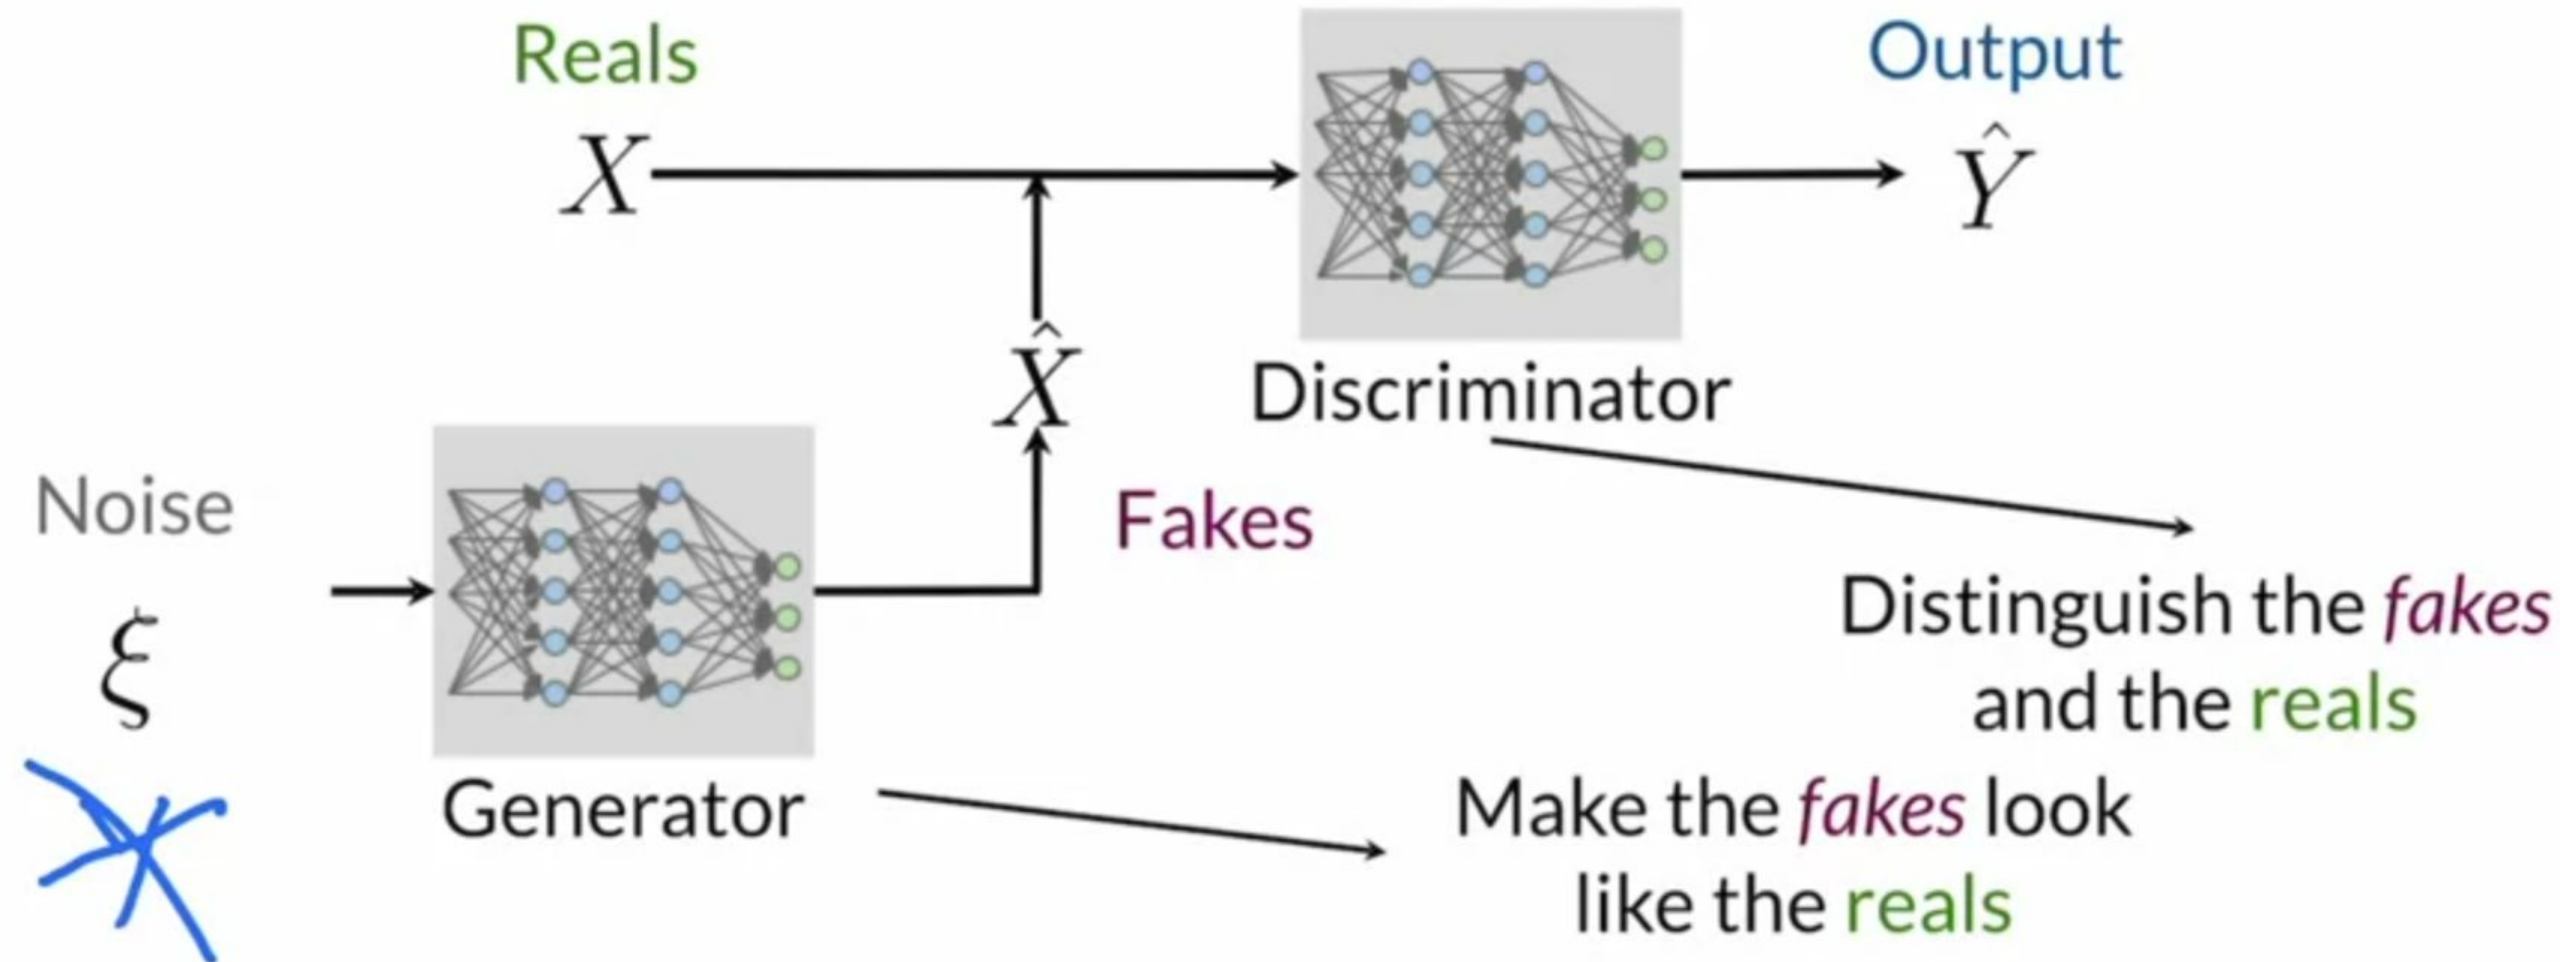
Ground truth is 0
but prediction is 1,
loss is ∞

Loss is ~ 0


- In summary, the BCE cost function has two main terms that are relevant for each of the classes
 - When the prediction and the label are similar, the BCE loss is close to 0
 - When they are very different, BCE loss approaches infinity
- The BCE loss is performed across a mini-batch of several examples, say n examples
 - It takes the **average** of all the examples
 - Each of those examples can be different. One of them can be 1, the other four could be 0, for their different classes.

Putting all together

GANs Model

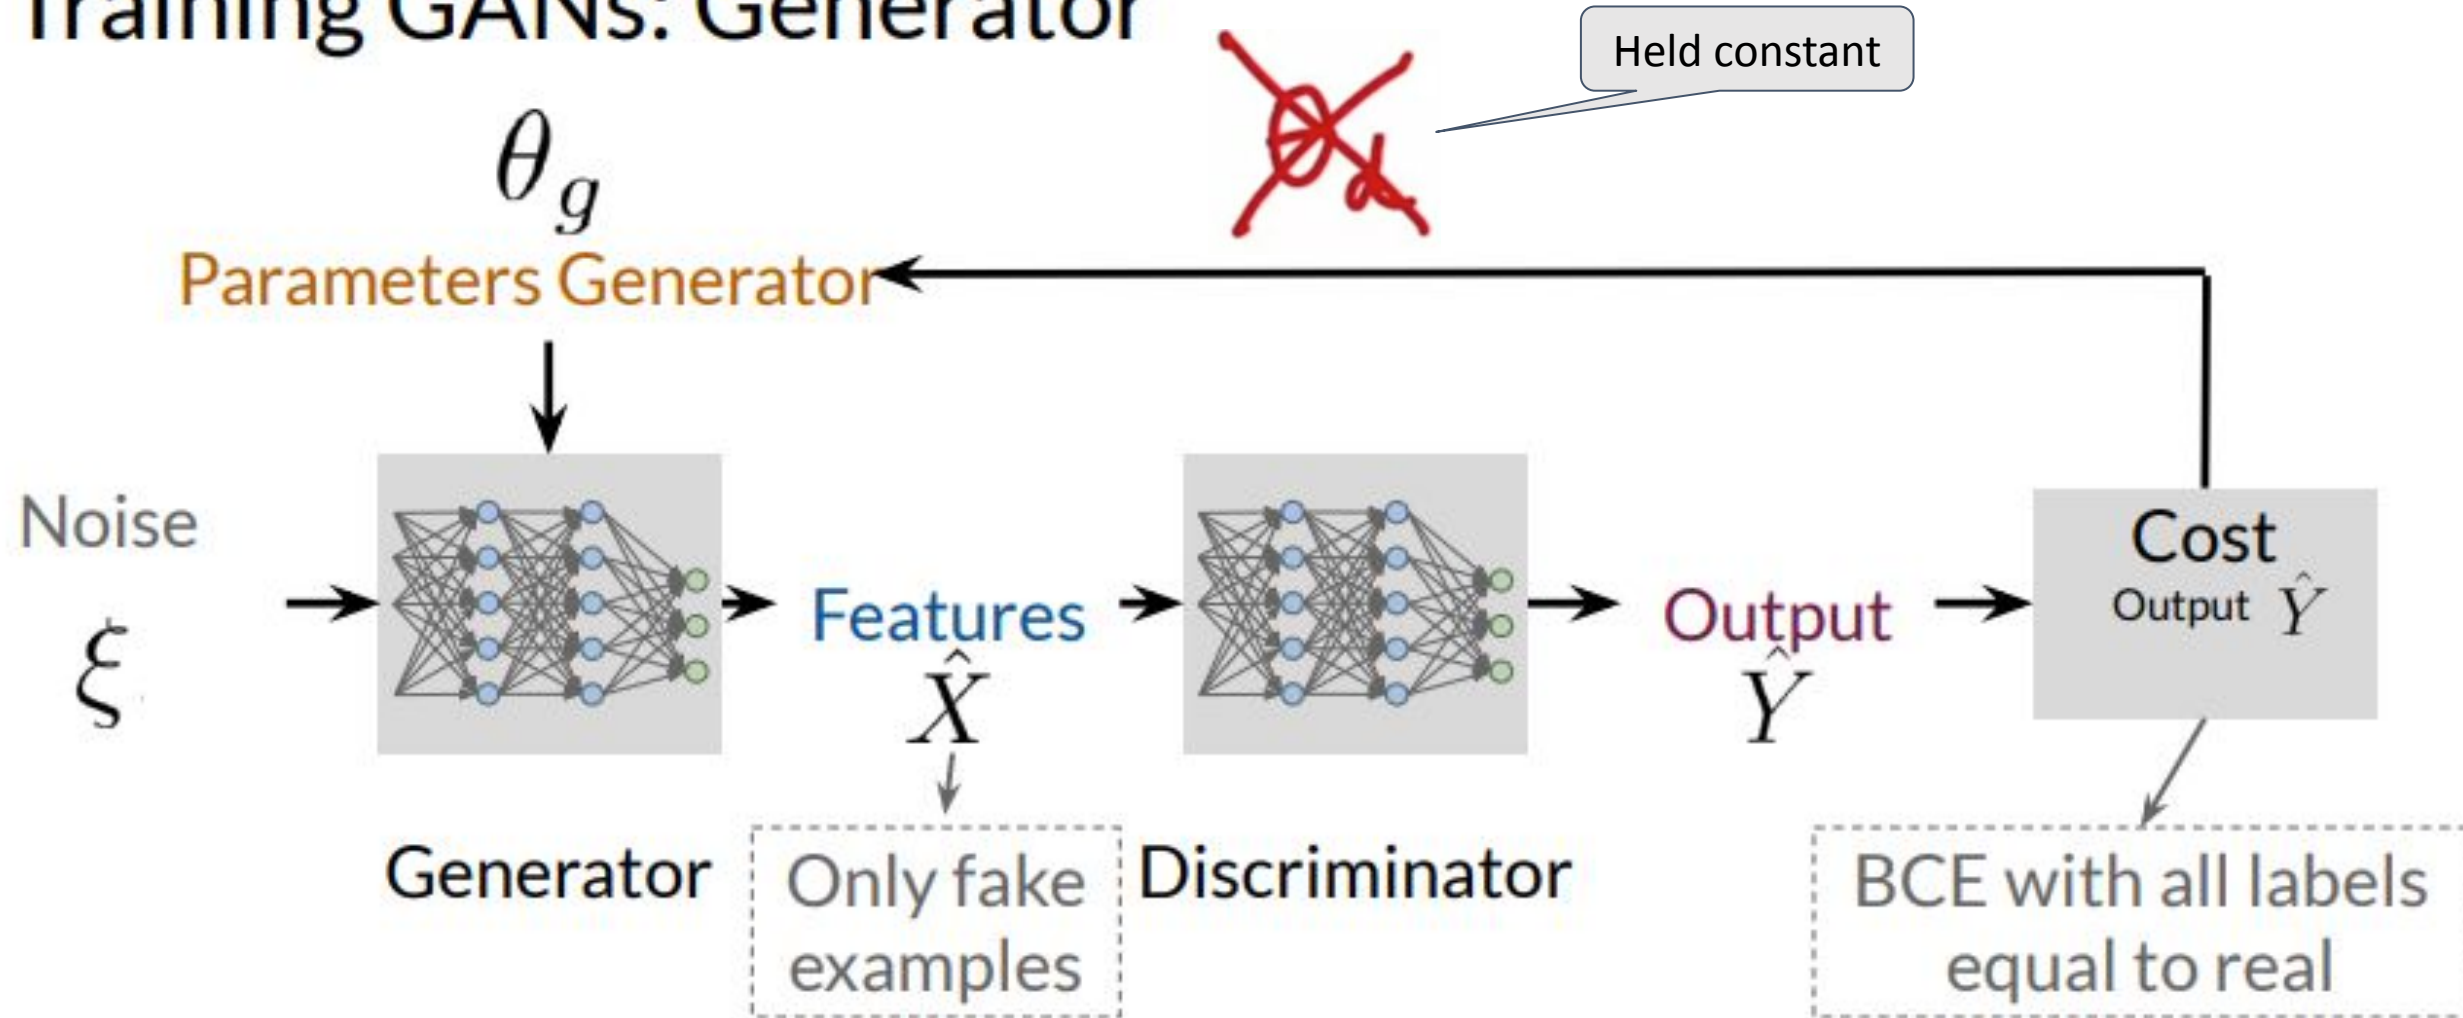


Training GANs: Discriminator



The two models are trained alternatively

Training GANs: Generator



Keep both models at the same level

- GANs train in alternating fashion
- The two models should always be at a similar “skill” level
 - If one model significantly better than the other, it doesn’t help the other learn because the feedback is not useful.
 - Imagine if you were a beginning artist, and you showed your work to an art expert, asking whether your painting looked like a famous piece and all they said was ‘no’. Because they have a very discerning eye, they know your image is not right, but won’t be able to tell you how close you are.

Training GANs



Handson Session-1

GAN